

```
import os
import json

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_squared_error

# Beautiful global style
sns.set_theme(style="whitegrid", palette="viridis")

# -----
# 1. FILE PATHS
# -----

FILE_PATHS = {
    2019: "/Users/randalburks/Downloads/FALL 2025/482/final
project/play_by_play_2019.csv",
    2020: "/Users/randalburks/Downloads/FALL 2025/482/final
project/play_by_play_2020.csv",
    2021: "/Users/randalburks/Downloads/FALL 2025/482/final
project/play_by_play_2021.csv",
    2023: "/Users/randalburks/Downloads/FALL 2025/482/final
project/play_by_play_2023.csv",
}

BASE_DIR = os.path.dirname(next(iter(FILE_PATHS.values())))

# -----
# 2. LOAD & CLEAN DATA
# -----

def load_pbp_data(file_paths):
    tables = []
    for season, path in file_paths.items():
        print(f"Loading season {season} from {path} ...")
        df = pd.read_csv(path, low_memory=False)
        df["season"] = season
        tables.append(df)
    full_df = pd.concat(tables, ignore_index=True)
    print("Total rows:", len(full_df))
    return full_df
```

```

def basic_cleaning(df):
    rename_map = {}
    if "total_home_score" in df and "home_score" not in df:
        rename_map["total_home_score"] = "home_score"
    if "total_away_score" in df and "away_score" not in df:
        rename_map["total_away_score"] = "away_score"

    df = df.rename(rename_map)

    keep_cols = [
        "season", "game_id", "qtr", "game_seconds_remaining", "score_differential",
        "posteam", "defteam", "home_team", "away_team",
        "home_score", "away_score", "epa", "wpa"
    ]

    df = df[keep_cols].dropna()
    print("Rows after cleaning:", len(df))
    return df

# -----
# 3. CLUTCH LABELING
# -----

def add_clutch_flag(df):
    df["is_clutch"] = (
        (df["qtr"] == 4) &
        (df["game_seconds_remaining"] <= 5*60) &
        (df["score_differential"].abs() <= 8)
    )
    return df

def get_clutch_plays(df):
    clutch = df[df["is_clutch"] == True].copy()
    print("Clutch plays:", len(clutch))
    return clutch

# -----
# 4. TEAM CLUTCH METRICS
# -----

def build_team_clutch_stats(clutch_df):
    clutch_df["success"] = clutch_df["epa"] > 0

```

```

grouped = clutch_df.groupby(["postteam", "season"]).agg(
    clutch_plays_count=("game_id", "count"),
    clutch_epa_avg=("epa", "mean"),
    clutch_wpa_avg=("wpa", "mean"),
    clutch_success_rate=("success", "mean")
).reset_index().rename(columns={"postteam": "team"})

print("Team clutch stats:", grouped.shape)
return grouped

# -----
# 5. TEAM WINS
# -----

def compute_team_wins(df):
    last_play =
df.sort_values(["game_id", "game_seconds_remaining"]).groupby("game_id"
).tail(1)

    def choose_winner(row):
        if row["home_score"] > row["away_score"]:
            return row["home_team"]
        elif row["away_score"] > row["home_score"]:
            return row["away_team"]
        return None

    last_play["winner"] = last_play.apply(choose_winner, axis=1)

    wins =
last_play.dropna(subset=["winner"]).groupby(["winner", "season"]).size(
).reset_index(name="wins")
    wins = wins.rename(columns={"winner": "team"})
    print("Team wins:", wins.shape)
    return wins

# -----
# 6. BASE MODEL – Wins ~ Clutch EPA
# -----

def wins_model(team_clutch, team_wins):
    merged = pd.merge(team_clutch, team_wins, on=["team", "season"],
how="left").fillna({"wins": 0})

    X = merged[["clutch_epa_avg"]]
    y = merged["wins"]

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

print("\n==== MODEL: Wins ~ Clutch EPA ===")
print("R2:", r2_score(y_test, y_pred))
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("RMSE:", rmse)
print("Slope:", model.coef_[0])
print("Intercept:", model.intercept_)

return merged, model

# -----
# 7. ADVANCED PREDICTIVE MODEL
# -----


def predictive_clutch_model(merged_stats, pbp_clean):
    overall_epa = pbp_clean.groupby("postteam")["epa"].mean()
    merged_stats["overall_epa"] =
    merged_stats["team"].map(overall_epa)

    feature_cols = [
        "clutch_success_rate",
        "clutch_plays_count",
        "clutch_wpa_avg",
        "wins",
        "overall_epa"
    ]

    X = merged_stats[feature_cols]
    y = merged_stats["clutch_epa_avg"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.3, random_state=42
    )

    model = RandomForestRegressor(n_estimators=300, random_state=42)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    print("\n==== ADVANCED PREDICTION MODEL ===")
    print("Predicting Clutch EPA")
    print("R2:", r2_score(y_test, y_pred))
    print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))

    importance_df = pd.DataFrame({

```

```

    "Feature": feature_cols,
    "Importance": model.feature_importances_
}).sort_values("Importance", ascending=False)

print("\n==== FEATURE IMPORTANCE ===")
display(importance_df)

latest = merged_stats["season"].max()
latest_df = merged_stats[merged_stats["season"] == latest].copy()

latest_df["predicted_clutch_epa"] =
model.predict(latest_df[feature_cols])

predicted_best = latest_df.sort_values("predicted_clutch_epa",
ascending=False).iloc[0]

print("\n==== PREDICTED BEST CLUTCH TEAM ===")
print("Team:", predicted_best["team"])
print("Predicted Clutch EPA:",
predicted_best["predicted_clutch_epa"])

return latest_df, model

# -----
# 8. NO-SQL JSON OUTPUT
# -----


def save_team_results_json(df):
    results = {}

    for _, row in df.iterrows():
        key = f"{row['team']}_{int(row['season'])}"
        results[key] = {
            "clutch_plays_count": int(row["clutch_plays_count"]),
            "clutch_epa_avg": float(row["clutch_epa_avg"]),
            "clutch_wpa_avg": float(row["clutch_wpa_avg"]),
            "clutch_success_rate": float(row["clutch_success_rate"]),
            "wins": int(row["wins"])
        }

    out_path = os.path.join(BASE_DIR, "team_clutch_stats.json")
    try:
        with open(out_path, "w") as f:
            json.dump(results, f, indent=2)
        print("Saved JSON.")
    except:
        print("JSON save skipped.")

```

```

# -----
# 9. BEAUTIFUL STORY PLOTS
# -----

def make_story_plots(df, clutch_df, merged_stats):

    # EPA distribution clutch vs non
    non_clutch = df[~df["is_clutch"]]
    non_sample = non_clutch.sample(min(len(non_clutch),
len(clutch_df)*3), random_state=1)

    plot_df = pd.concat([
        clutch_df.assign(situation="Clutch")[["epa", "situation"]],
        non_sample.assign(situation="Non-Clutch")[["epa", "situation"]]
    ])

    plt.figure(figsize=(10,7))
    sns.histplot(plot_df, x="epa", hue="situation", bins=50,
stat="density", alpha=0.5)
    plt.xlim(-5,5)
    plt.title("EPA Distribution – Clutch vs Non-Clutch")
    plt.show()

    # Team ranking
    latest = merged_stats["season"].max()
    latest_df =
merged_stats[merged_stats["season"]==latest].sort_values("clutch_epa_avg")

    plt.figure(figsize=(10,8))
    sns.barplot(data=latest_df, x="clutch_epa_avg", y="team")
    plt.title(f"Clutch EPA by Team (Season {latest})")
    plt.show()

    # Clutch EPA vs Wins
    plt.figure(figsize=(8,6))
    sns.regplot(data=merged_stats, x="clutch_epa_avg", y="wins",
scatter_kws={"alpha":0.7})
    plt.title("Clutch EPA vs Wins")
    plt.show()

# -----
# 10. EDA + IDA
# -----


def eda_overview(df):
    print("===== HEAD =====")

```

```

display(df.head())
print("\n===== SUMMARY =====")
display(df.describe())
print("\n===== MISSING VALUES =====")
display(df.isnull().sum())

def ida_clutch_vs_nonclutch(df):
    clutch = df[df["is_clutch"]]
    non = df[~df["is_clutch"]]
    summary = pd.DataFrame({
        "Mean EPA": [clutch["epa"].mean(), non["epa"].mean()],
        "Std EPA": [clutch["epa"].std(), non["epa"].std()],
        "Count": [len(clutch), len(non)]})
    , index=["Clutch", "Non-Clutch"])

    print("===== CLUTCH VS NON-CLUTCH =====")
    display(summary)

def ida_team_level(merged_stats):
    pivot = merged_stats.pivot(index="team", columns="season",
values="clutch_epa_avg")

    plt.figure(figsize=(12,9))
    sns.heatmap(pivot, cmap="magma", center=0)
    plt.title("Clutch EPA Heatmap by Team & Season")
    plt.show()

# -----
# 11. CLUTCH IMPROVEMENT SCORE
# -----


def plot_clutch_improvement(df, merged_stats):
    overall = df.groupby("posteam")["epa"].mean()
    merged_stats["overall_epa"] = merged_stats["team"].map(overall)
    merged_stats["clutch_improvement"] =
    merged_stats["clutch_epa_avg"] - merged_stats["overall_epa"]

    plt.figure(figsize=(10,8))
    sns.barplot(
        data=merged_stats.sort_values("clutch_improvement"),
        x="clutch_improvement", y="team", palette="viridis"
    )
    plt.title("Clutch EPA Improvement (Clutch EPA - Overall EPA)")
    plt.show()

```




```
print("\nPredicted Best Clutch Team:", predicted_best)


# RUN EVERYTHING
main()


Loading season 2019 from /Users/randalburks/Downloads/FALL
2025/482/final project/play_by_play_2019.csv ...
Loading season 2020 from /Users/randalburks/Downloads/FALL
2025/482/final project/play_by_play_2020.csv ...
Loading season 2021 from /Users/randalburks/Downloads/FALL
2025/482/final project/play_by_play_2021.csv ...
Loading season 2023 from /Users/randalburks/Downloads/FALL
2025/482/final project/play_by_play_2023.csv ...
Total rows: 194552
Rows after cleaning: 181295
Clutch plays: 9894
Team clutch stats: (128, 6)
Team wins: (128, 3)

==== MODEL: Wins ~ Clutch EPA ===
R2: 0.017253156547390236
RMSE: 3.286611588293964
Slope: 9.501049206676011
Intercept: 8.872312359096965
Saved JSON.
```



Clutch EPA by Team (Season 2023)





```
===== HEAD =====
```

	season	game_id	qtr	game_seconds_remaining
0.0	2019	2019_01_ATL_MIN	1	3600.0
0.0	2019	2019_01_ATL_MIN	1	3600.0
0.0	2019	2019_01_ATL_MIN	1	3560.0
0.0	2019	2019_01_ATL_MIN	1	3521.0
0.0	2019	2019_01_ATL_MIN	1	3479.0

	posteam	defteam	home_team	away_team	home_score	away_score
0.000000	ATL	MIN	MIN	ATL	28	12
2	ATL	MIN	MIN	ATL	28	12 -

```
1.658763
3   ATL     MIN     MIN     ATL      28      12  -
0.538914
4   ATL     MIN     MIN     ATL      28      12
0.142138
5   ATL     MIN     MIN     ATL      28      12  -
4.034299
```

```
          wpa  is_clutch
1  0.000000    False
2 -0.035015    False
3 -0.018435    False
4  0.018007    False
5 -0.113866    False
```

===== SUMMARY =====

```
          season        qtr game_seconds_remaining \
count  181295.000000  181295.000000           181295.00000
mean   2020.779122    2.540434            1745.19006
std    1.479901       1.124586            1041.55892
min   2019.000000     1.000000             0.00000
25%   2020.000000     2.000000            855.00000
50%   2021.000000     2.000000            1803.00000
75%   2023.000000     4.000000            2641.00000
max   2023.000000     5.000000            3600.00000
```

```
          score_differential    home_score    away_score      epa
\count  181295.000000  181295.000000  181295.000000  181295.000000
mean    -1.459902      23.959833     22.736330    -0.002938
std     10.696310      10.132413     10.052614     1.295115
min    -56.000000      0.000000     0.000000    -13.584859
25%   -7.000000       17.000000     16.000000    -0.581829
50%   0.000000       24.000000     23.000000    -0.034859
75%   4.000000       30.000000     30.000000     0.590608
max   55.000000      70.000000     59.000000     8.882384
```

```
          wpa
count  181295.000000
mean   0.000845
std    0.040171
```


```
min      -0.737291
25%     -0.014271
50%     -0.000019
75%      0.011980
max      0.944656
```

===== MISSING VALUES =====

```
season          0
game_id         0
qtr             0
game_seconds_remaining 0
score_differential 0
posteam          0
defteam          0
home_team        0
away_team        0
home_score       0
away_score       0
epa              0
wpa              0
is_clutch        0
dtype: int64
```

===== CLUTCH VS NON-CLUTCH =====


	Mean EPA	Std EPA	Count
Clutch	-0.014367	1.245143	9894
Non-Clutch	-0.002279	1.297942	171401



```
/var/folders/j9/y9drl51s2gvbm203cbmk7py80000gn/T/
ipykernel_71478/2222944035.py:324: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `y` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.barplot(
```



===== BEST CLUTCH TEAM (REAL DATA) =====

Team: SEA
 Clutch EPA: 0.1619354368476224
 Success Rate: 0.4807692307692308

Real Best Clutch Team: SEA

==== ADVANCED PREDICTION MODEL ===

Predicting Clutch EPA
 R^2 : 0.8126171716994586
 RMSE: 0.07263050895451344

==== FEATURE IMPORTANCE ===

	Feature	Importance
2	clutch_wpa_avg	0.754805
0	clutch_success_rate	0.139277
1	clutch_plays_count	0.046034

```
4      overall_epa    0.039720
3          wins     0.020164
```

```
==== PREDICTED BEST CLUTCH TEAM ====
```

```
Team: SEA
```

```
Predicted Clutch EPA: 0.12703221419364158
```

```
Predicted Best Clutch Team: SEA
```